

SHREYASH GONDANE

Staff Platform Engineer — MCP · Auth & Identity · API Infrastructure · Distributed Systems

(+1) 929-527-9683 | sg6874@nyu.edu | [Portfolio](#) | [LinkedIn](#) | [GitHub](#) | [LeetCode](#)

PROFESSIONAL SUMMARY

Platform engineer with 4+ years of compounded scope — leading architecture, owning end-to-end technical vision, and shipping foundational platform primitives that other engineers depend on daily. Built and maintained production MCP integrations and A2A Protocol infrastructure: the exact connectivity layer Anthropic's Platform Connectivity team owns. Designed enterprise-grade Auth & Identity systems (Auth0, Okta, Cognito — JWT, SAML, MFA, RBAC, SSO) at $\leq 3\text{ms}$ token latency with zero CVEs — a production Auth platform. Shipped the complete API Distributability stack: multi-cloud Terraform (AWS/GCP/Azure), cloud-agnostic deployment pipelines, and enterprise compliance infrastructure. Contributed directly to Scale AI's LLM Engine: enterprise Prometheus monitoring, priority-based batch orchestration, and CI/CD infrastructure — giving hands-on familiarity with AI platform architecture at production scale. Track record of 0-to-1 platform builds, cross-functional technical leadership, and architectural decisions with lasting impact. M.S. Computer Science, NYU. 2 Springer publications.

TECHNICAL SKILLS

Languages: Python (primary, 5+ yrs), TypeScript / Node.js, Go (systems tooling), Java 11, SQL, Bash, Rust (familiar)

Platform — Connectivity & MCP: MCP proxy design, A2A Protocol orchestration, OAuth2 / OIDC / token refresh at scale, API gateway architecture, multi-tenant platform engineering, enterprise SDK authoring (PyPI, npm), service mesh

Auth & Identity: Auth0, Okta, AWS Cognito — JWT, SAML, MFA, RBAC, SSO, row-level security; KMS envelope encryption; IAM least-privilege; zero CVE track record at $\leq 3\text{ms}$ token latency

API Distributability: Multi-cloud Terraform (AWS/GCP/Azure), cross-cloud networking, EKS, Helm, Docker, Kubernetes, cloud-agnostic deployment pipelines, enterprise compliance (SOC-2, SEC 17a-4, FedRAMP-adjacent)

Platform Acceleration: OpenTelemetry (traces, metrics, logs), Prometheus, Grafana, Datadog, CloudWatch, Sentry, CI/CD automation (GitHub Actions, Jenkins), developer tooling, load testing, SLO/SLA design

Platform Intelligence: RAG pipelines (Vector + Vectorless), LLM serving infrastructure (Claude / GPT-4 / Titan / Llama / Bedrock), model eval harnesses, RHFL fine-tuning loops, LangChain, LangGraph, embedding pipelines, token cost governance

Data & Messaging: PostgreSQL (RLS, vector search, query optimization), DynamoDB, Redis, Kafka, OpenSearch, Aurora, SQS, SNS

INDUSTRY EXPERIENCE

Staff-Level Platform Engineer — **Engineersmind Corp.** — **Jersey City, NJ** Jan 2024 – Present

AI Platform — Foundational Primitives & Infrastructure ([pdffillr.ai](#)) | Lead Platform Architect

Python · Pydantic · MCP · A2A Protocol · LangChain · LangGraph · Lambda · Step Functions · SQS · DynamoDB · OpenSearch · KMS · Terraform (AWS/GCP/Azure) · OTel · Prometheus · Grafana · Jenkins

- **MCP & Connectivity Infrastructure** — Designed and deployed production Model Context Protocol (MCP) integrations and Agent-to-Agent (A2A) communication layer connecting Classification, Extraction, Mapping, and Validation agents — the exact infrastructure Anthropic's Connectivity team owns. Built the proxy routing layer, session management, and fallback patterns for agent-to-service communication with enterprise reliability.
- **API Distributability — Multi-Cloud Architecture** — Authored Terraform modules for AWS, GCP, and Azure enabling cloud-agnostic deployment of the full platform stack (compute, storage, queues, secrets, monitoring); local development parity via Docker Compose; CI/CD-ready with drift detection. Reduced document preparation from 15 days → 5 minutes at 97% field accuracy — a 0-to-1 platform build with lasting impact.
- **Platform Acceleration — CI/CD & Observability** — Built automated CI/CD pipelines with security scanning on every commit; OTel distributed tracing + Prometheus + Grafana for full platform observability; SLO enforcement and cost governance dashboards cut compute cost ~50%. Published SDKs to PyPI and npm with versioned documentation — the developer experience layer that multiplies engineer productivity.
- **Platform Intelligence — ML Training & Adaptation Loop** — Implemented Reinforcement Human Feedback Loop (RHFL) with Researcher-in-the-Loop with VLLM inference updates cycles for continuous model accuracy improvement (87% auto-fill, ~90% → 97% field accuracy); multi-model orchestration (Claude 3.5, GPT-4o, Llama 3.2, HuggingFace distilled models) with per-request token consumption tracking and cost attribution — production ML pipeline with task-specific adaptation at its core.
- **Architectural Ownership & Cross-Functional Leadership** — Independently scoped and drove multi-month architectural RFCs; set technical standards for the platform; mentored 4 engineers on Clean Architecture and secure coding; raised PR acceptance rate 30%. Made build-vs-buy decisions, evaluated third-party integrations, and owned platform health end-to-end from design through production operations.

- **Auth & Identity Platform** — Built Anthropic's Auth & Identity problem from scratch: multi-tenant Auth0/Okta + AWS Cognito (JWT, SAML, MFA, RBAC, SSO) at ≤3ms token latency with zero CVEs across 6 microservices. PostgreSQL RLS with tenant and role claims for data isolation. Cloudflare Workers for HS256/RS256 JWT verification at the edge. This is the production identity infrastructure Anthropic's Auth team would own.
- **Connectivity & API Gateway** — Built geo-sharded routing and request normalization layer via Cloudflare Workers (TypeScript); Logpush → S3 → OTEL Collector for full trace/metric coverage; Webhooks → SQS → Lambda → PostgreSQL pipeline at 18K events/hour sub-second latency; Redis read-through caching cut API latency 27% — proxy infrastructure designed for reliability when downstream services don't stay up.
- **Platform Acceleration — Microservices & Developer Productivity** — Independently designed and executed the split of a monolith into 6 production microservices (Nx, EKS, Docker); OTEL + Grafana Tempo distributed tracing; SLO dashboards achieved 2× peak traffic with zero downtime. Drove architectural RFCs and made technology selection decisions independently, operating as the technical authority for the platform.
- **LLM Infrastructure & Platform Intelligence** — Operated GPT-4 production inference (1.2M conversations/month at p95 < 350ms); built GPT-4 vs. Claude vs. Titan benchmarking harness with per-request cost governance; RAG pipeline (OpenAI embeddings + DynamoDB cache) with hallucination guardrails reduced errors 42% on QA evaluation sets. Led team of 21 engineers.

Platform Architect — Engineersmind — Hedge Funds Portal & FINRA Compliance Archival2021 – 2024

- **Enterprise Security & Compliance Platform** — Designed multi-layer PGP + KMS envelope encryption with role-segregated IAM least-privilege; SOC-2 & SEC 17a-4 compliant; passed 2 external penetration tests without findings. Built OAuth 2.0 client + webhook handler driving ~7K digital-sign events/week; reduced signature cycle 48h → 6h. Enterprise compliance posture directly applicable to Anthropic's API Distributability security requirements.
- **Large-Scale Platform Reliability** — Serverless ingestion framework (SFTP + SMTP → Lambda → S3 → DynamoDB) processing 200 GB/year at 45 MB/s; 20 Lambda alert rules with SecOps cut false-positives 35%; CloudWatch → Grafana drove 42% MTTR reduction with 99.9% SLO uptime over 2 years; zero data-loss incidents.

System Engineer — Tata Consultancy Services — MH, IndiaFeb 2020 – Jul 2021

- **Production Systems at Scale** — Banking APIs serving 2,000+ branches and 15M customers; 30% query performance gain via index optimization; 50% efficiency improvement through caching and server migration; OCI Monitoring + Prometheus for real-time reliability.

OPEN SOURCE — DIRECT ANTHROPIC ECOSYSTEM CONTRIBUTIONS

- **Scale AI LLM Engine — Platform Infrastructure (2025):** Directly extended Scale AI's production LLM Engine: enterprise Prometheus metrics integration, priority-based batch job scheduler with exponential backoff retry, mock executor enabling CPU-only testing (zero GPU cost), and local-to-cloud portable deployment infrastructure. 50× faster scheduling (5s → 0.1s); 100% test coverage, 85 tests, 37s CI runtime. Deep working knowledge of LLM platform architecture at production scale — adjacent to Anthropic's own Platform teams.
- **Terraform IaC — Multi-Cloud Platform System (2025):** Production-grade AWS infrastructure (EC2, RDS PostgreSQL, S3, VPC isolation, Secrets Manager, IAM least-privilege, remote state S3 + DynamoDB locking) with CI/CD-ready Terraform workflows and drift detection — the cloud-agnostic tooling pattern Anthropic's Multicloud team builds.
- **Multi-Agent Trading System — Platform Decision Architecture (2025):** 7 specialized AI agents with consensus-based signal aggregation; backtested 15 symbols / 10 years with Monte Carlo simulation (10,000 paths) — demonstrating multi-agent platform architecture at scale.
- **Open-DLA — ML Evaluation Platform (NYU 2023–Present):** Unified ML benchmarking framework (Factory Method pattern); open-sourced with integration tests, versioned releases, and documentation — developer experience and platform tooling that engineers love to use.

RESEARCH & EDUCATION

New York University — M.S. Computer Science (Distributed Systems, ML Infrastructure, Security & Platform Engineering)

Publications: Virtual Assistant Chatbot Architecture based on NLP for Customer Service — RAAI 2020, Springer • A Machine Learning Based Approach to Password Authentication using Keystroke Biometrics — Springer AI Conference